

Clean → Mary

When Mary and John went to the store, **John** gave a drink to

Corrupted → John

When Mary and John went to the store, **Mary** gave a drink to

1

Run the clean prompt and store every internal activation

2

Run the corrupted prompt, but overwrite ONE activation (one layer, one position, or one head) with its clean value

3

Measure how much of the clean answer comes back

$\text{logit diff} = \text{logit}(\text{Mary}) - \text{logit}(\text{John})$

$\text{score} = (\text{patched} - \text{corrupted}) / (\text{clean} - \text{corrupted})$

1 = answer fully restored · 0 = no effect